



CMP 28

Household Electricity

Submitted by
Supercaliprobablisticexpialidocious

Our Beloved Team

Name	Code
مصطفى محمد عبداللاه محمد	91240765
مسلم احمد محمد اسماعيل سيد	91241062
احمد ماجد محمد خطابي السيد خطابي	91240928
عبدالرحمن جلال شعبان السيد	91240413
على ثروت على احمد محمود	91240474
عبدالله طارق سيد احمد محمود	91240462

Presented to:

Prof. Maha Hassanein

Academic Year: 2025/2026

Abstract

This research project investigates individual household electricity consumption patterns using probability theory and statistical inference methods. Analyzing **2,075,259 complete observations** from a household in Sceaux, France (December 2006 - November 2010), we employed **+10 statistical tests** across parametric, non-parametric, and distribution-fitting frameworks to comprehensively characterize consumption distributions and identify temporal patterns.

Key Findings: (1) Global active power consumption follows a **Weibull distribution** with exceptional fit quality, superior to Normal, and Lognormal; (2) **Significant weekend effect:** weekends show higher consumption; (3) **Pronounced seasonal variation:** Consumption in Winter is higher than Summer; (4) **Strong linear relationship:** Global intensity predicts power variance; (5) **Significant categorical associations:** Season $\times$ Consumption and Hour $\times$ Consumption. These findings provide empirical foundation for smart grid optimization, demand-response program design, and infrastructure planning, with practical implications for capacity expansion and renewable integration in temperate-climate regions.

1 Introduction

1.1 Background

Household electricity consumption has strong implications for utility companies, grid operators, policy makers, consumers, and researchers. Traditional analysis shows that aggregate consumption across many homes obscure the behavioral and temporal details needed for smart-grid applications. Single-household analysis reveals daily cycles, seasonal shifts, occupancy-dependent patterns, and appliance-level signatures that are essential for accurate short-term forecasting, peak management, demand-response design, and probabilistic risk assessment. The project therefore asks whether individual (hourly) household consumption follows identifiable probability distributions that can be used for prediction and extreme-event assessment across multiple temporal scales.

1.2 Objectives

1.2.1 Primary study objective:

- 1.2.2 Assess whether realistic synthetic hourly electricity-consumption data can be generated in a way that preserves the statistical characteristics of the observed real-world consumption.

1.2.3 Specific study objectives:

1. Characterize the probability distribution of household electricity consumption.
2. Identify and quantify temporal patterns at hourly, daily and seasonal scales.
3. Test for day-of-week effects (weekday vs weekend).
4. Analyze seasonal variations (winter, spring, summer, fall).
5. Model linear relationships between electrical variables and consumption.
6. Generate synthetic hourly consumption data that preserve observed statistical properties.
7. Assess probabilities and return periods of extreme consumption events (e.g., >95th percentile).

8. Test independence/association between categorical variables (hour, day type, season) and consumption levels.

1.3 Significance

- Enables utilities and grid operators to improve 24-48 hour load forecasts, optimize generation mix, and reduce risk of overloads and blackouts.
- Supports peak demand management and realistic target setting for demand-response programs by quantifying consumption variability and extremes.
- Informs policy makers for targeted energy-conservation and renewable-integration policies based on household-level behavior.
- Empowers consumers with personalized energy-management insights and cost-reduction strategies.
- Provides researchers with validated distributional models, synthetic datasets, and methods for anomaly detection, probabilistic forecasting, and extreme-event risk assessment.

1.4 Research Questions and Statistical Hypotheses

1.4.1 RQ1: Distribution Characterization

- **Research Question:** Which probability distribution (Normal, Lognormal, Gamma, Weibull) best characterizes household electricity consumption?
- **Hypotheses:**
 - H_0 : Data follow a normal distribution
 - H_a : Data do not follow a normal distribution

1.4.2 RQ2: Day-Type Effects (Weekday vs. Weekend)

- **Research Question:** Is consumption significantly different on weekends versus weekdays?
- **Hypotheses:**
 - $H_0: \mu_{weekday} = \mu_{weekend}$ (equal means)
 - $H_a: \mu_{weekday} \neq \mu_{weekend}$

1.4.3 RQ3: Seasonal Effects

- **Research Question:** Do seasonal variations significantly affect consumption patterns?
- **Hypotheses:**
 - $H_0: \mu_{summer} = \mu_{spring} = \mu_{fall} = \mu_{winter}$
 - H_a : At least one seasonal mean differs significantly

1.4.4 RQ4: Predictability and Linear Relationships

- **Research Question:** Can we accurately predict consumption from electrical variables?
- **Hypotheses:**
 - $H_0: \beta_1 = 0$ (no linear relationship between intensity and power)
 - $H_a: \beta_1 \neq 0$ (significant relationship exists)

1.4.5 RQ5: Independence of Categorical Variables

- **Research Question:** Are categorical variables (hour, day type, season) independent of consumption levels?

- **Hypotheses:**
 - H_0 : Categorical variable and consumption level are independent
 - H_a : Variables are dependent (associated)

1.5 Response Variables and Factors

1.5.1 Primary Response Variable:

- **Global Active Power (kilowatts):** The total active power consumed by the household, measured at one-minute intervals. This represents instantaneous electrical load.

1.5.2 Secondary Response Variables:

- **Global Reactive Power (kilowatts):** Reactive power component of total consumption
- **Voltage (volts):** Average voltage measured at the household's connection
- **Global Intensity (amperes):** Average current intensity
- **Sub-metering values (watt-hours):** Circuit-specific consumption for:
 - **(Sub_metering_1):** Kitchen
 - **(Sub_metering_2):** Laundry
 - **(Sub_metering_3):** Climate Control

1.5.3 Controlled Factors:

- Time of day (hour of day, 0-23)
- Day type (weekday vs. weekend)
- Season (winter, spring, summer, fall)
- Date (specific day from Dec 2006 - Nov 2010)

1.5.4 Uncontrolled Factors:

- Occupancy patterns (number of people present, presence duration)
- Ambient temperature (affects heating/cooling loads)
- Weather conditions (humidity, solar radiation affecting cooling demand)
- Appliance usage decisions (user behavior, lifestyle choices)
- Appliance efficiency (age and condition of electrical equipment)
- Seasonal activities (vacation periods, extreme weather)

1.6 Data and Variables

1.6.1 Data Source

The dataset is obtained from the UCI Machine Learning Repository, contributed by Georges Hebrail and Alice Berard (Hebrail & Berard, 2012). The data represents measurements from a single household located in Sceaux, France (7 km from Paris), collected between December 2006 and November 2010 (47 months). Measurements were recorded at one-minute sampling intervals, yielding 2,075,259 observations with nine variables.

Data Quality Characteristics:

- Missing values: Approximately 0.97% of rows contain missing measurement values
- License: Creative Commons Attribution 4.0 International (CC BY 4.0)
- Accessibility: Free download from <https://archive.ics.uci.edu/ml/datasets/individual+household+electric+power+consumption>

1.6.2 Variables to analyze


Variable Name	Type	Units	Description	Expected Range
Date	Date	--	Calendar date of measurement	16/12/2006 26/11/2010
Time	Time	--	Time of day	00:00:00 – 23:59:59
Global Active Power	Continuous	Kilowatt (kW)	Household total active power, minute averaged	0.076 - 11.122 kW
Global Active Power	Continuous	Kilowatt (kW)	Household total reactive power, minute averaged	0.0 - 6.8 kW
Voltage	Continuous	Volt (V)	Average voltage over minute interval	223.0 - 254.0 V
Global Intensity	Continuous	Ampere (A)	Average current intensity over minute	0.2 - 48.4 A
Sub_metering 1	Continuous	Watt-hours (Wh)	Kitchen active energy (dishwasher, oven, microwave)	0 - 88 Wh
Sub_metering 2	Continuous	Watt-hours (Wh)	Laundry active energy (washer, dryer, fridge, light)	0 - 80 Wh
Sub_metering 3	Continuous	Watt-hours (Wh)	Climate control active energy (water heater, A/C)	0 - 31 Wh

1.7 Response Variables for Each Research Question

Research Question	Primary Response Variable	Expected Statistical Test
RQ1: Distribution Fit	Global_active_power	Goodness-of-Fit Tests (KS, Anderson-Darling, Chi-Square)
RQ2: Hourly Variation	Global_active_power	One-Way ANOVA, Kruskal-Wallis
RQ3: Weekday/Weekend	Global_active_power	Two-Way ANOVA with interaction
RQ4: Seasonal Effects	Global_active_power	One-Way ANOVA, post-hoc Tukey HSD
RQ5: Predictability and Linear Relationships	Global_active_power	Simple/Multiple Regression
RQ6: Independence of Categorical Variables	Global_active_power	Chi-Square of Independence


1.8 Descriptive Statistics

1.8.1 Observation Counts by Category




These bar charts show that the dataset is evenly distributed across seasons, hours, and day types, indicating no major sampling bias. Seasonal counts are nearly equal, hourly observations are consistently balanced across all 24 hours, and while weekdays have slightly more entries, weekends are still well represented. Overall, the dataset provides a solid, unbiased foundation for temporal analysis.

1.8.2 Hourly Consumption Pattern




The plots reveal a clear daily usage cycle, with low early-morning consumption, rising demand through the day, peak evening usage, and greater variability during peak hours.

1.8.3 Seasonal and Day Type Consumption Patterns




1.9 visualization shows higher and more variable electricity use on weekends and in Winter and Fall, reflecting the influence of daily routines and seasonal heating needs

Power Distribution Analysis




The eight-panel plots show that among Normal, Lognormal, Gamma, and Weibull distributions, the Weibull best fits household electricity consumption, closely matching both density and Q-Q plots.

1.10 Correlation between Active Power and Intensity



The graph shows that there's a very high correlation between Global Active Power and Global Intensity, which means that Global Intensity is **strong predictor** of consumption.

1.11 Analysis of Hourly Patterns



Beyond the aggregate means, we analyzed the hourly structure of demand to identify specific behavioral differences (in the figure above). We can see Hourly profiles showing distinct morning spikes on weekdays and sustained daytime loads on weekends, also the visual analysis reveals three distinct lifestyle signatures (mentioned in the results & analysis section).

2 Methodology

2.1 Statistical Tests: Purpose, Decision Criteria, and Applications

2.1.1 Test 1: Shapiro-Wilk Test (Normality)

What it does:

Checks whether sample data come from a normal distribution by comparing ordered sample values to expected normal quantiles. Particularly powerful for small to moderate samples.

Decision criterion:

- H_0 : Data are normally distributed
- If **p-value < 0.05** → **Reject H_0** : data are **NOT normal**
- If **p-value ≥ 0.05** → **Fail to reject H_0** : data **consistent with normality**

Test statistic: $W = \frac{[\sum_{i=1}^n a_i x_{(i)}]^2}{\sum_{i=1}^n (x_i - \bar{x})^2}$

- **Range:** 0→1
- **Interpretation:**
 - **Close to 1:** Perfect Bell Curve shape.
 - **Significantly < 1:** Skewed or distorted data.
 - *(Example: 0.79 indicates the data is NOT shaped like a bell)*

Limitations:

- **Small Data (N<50):** Very accurate.
- **Large Data (N>5000):** Too sensitive (flags tiny errors).
- **Solution:** For large data, run the test on a random sample (e.g., N=5000).

Metric	Value	Meaning
W	0.99	Very Normal
W	0.70	Very Skewed
p	< 0.05	NOT Normal (Significant)
p	> 0.05	Normal (Not Significant)

What questions it solves:

- "Can I assume normality to use parametric tests (t-test, ANOVA)?"
- "Does global active power follow a normal distribution?"
- "Is my data suitable for normal-theory statistical inference?"

2.1.2 Test 2: Anderson-Darling Test (Distribution Fit)

What it does:

Tests goodness-of-fit for a specified distribution, weighing deviations more heavily in the tails than Shapiro-Wilk. Sensitive to tail behavior.

Decision criterion:

- H_0 : Data follow the specified distribution
- If **Test Statistic > Critical Value** → **Reject H_0** (poor fit)
- If **Test Statistic ≤ Critical Value** → **Fail to reject H_0** (acceptable fit)

Test statistic: $A^2 = -n - \frac{1}{n} \sum_{i=1}^n (2i - 1) [\ln F(x_i) + \ln (1 - F(x_{n+1-i}))]$

Comparison	Meaning
Stat > Critical	NOT Normal (Significant deviation)
Stat < Critical	Normal (Fits the curve)
Larger Statistic	Worse fit (More skewed)
Smaller Statistic	Better fit (More normal)

- **Range:** $0 \rightarrow \infty$ (Positive numbers).
- **Interpretation:**
 - **Small Number:** The data fits the Bell Curve well.
 - **Large Number:** The data deviates from the Bell Curve (bad fit).
 - (Note: Unlike Shapiro, there is no 0-1 scale. You must compare it to Critical Values).

What questions it solves:

- "How well does the normal distribution fit my data?"
- "Are tail deviations significant or random?"

2.1.3 Test 3: Kolmogorov-Smirnov (KS) Test

What it does:

Compares the empirical cumulative distribution function (ECDF) of sample data to theoretical CDF of a specified distribution. Measures maximum vertical distance between curves.

Decision criterion:

- H_0 : Data follow the specified theoretical distribution
- If **p-value** < **0.05** → **Reject** H_0 (distribution does not fit)
- If **p-value** ≥ **0.05** → **Fail to reject** H_0 (fit is acceptable)
- In practice, for **very large datasets** (e.g., millions of rows), even tiny deviations make the p-value extremely small (often p=0), causing H_0 to be formally rejected.
- **Reality Check:** In large datasets, the p-value becomes unreliable. Instead, focus on the **KS Statistic (D)**. The distribution with the **lowest D** "wins" as it stays closest to the empirical data with minimal error.

Test statistic: $D = \max_x |F_n(x) - F(x)|$ where F_n = empirical CDF, F = theoretical CDF

- **Interpretation:**
 - **Small D:** The lines are tight. Good fit.
 - **Large D:** The lines are far apart. Bad fit.

The Verdict: "The Lowest Score Wins"

- When comparing multiple distributions (e.g., Normal vs. Weibull vs. Gamma), the distribution with the **lowest KS Statistic (D)** is the winner.
- **Why?** It means that specific shape stayed closest to your real data with the least amount of error.

The Reality Check: Why is p-value = 0?

- In large datasets (like ours with > 2 million rows), statistical tests become **extremely sensitive**.
- Even tiny, irrelevant imperfections make the test reject the fit mathematically (p<0.05).
- **Conclusion:** We ignore the p-value here and focus on the **KS Statistic (D)**. The lowest D wins, even if p=0.

What questions it solves:

- "Does consumption follow Normal/Lognormal/Gamma/Weibull distribution?"
- "Which candidate distribution provides best fit (smallest KS statistic)?"

2.1.4 Test 4: Chi-Square Goodness-of-Fit Test

What it does:

Compares observed frequencies in bins to expected frequencies under theoretical distribution. It's a non-parametric alternative to KS test suitable for categorical or binned data.

Decision criterion:

- H_0 : Observed frequencies follow theoretical distribution

- Compute: $\chi^2 = \sum_{i=1}^k \frac{(O_i - E_i)^2}{E_i}$ where O_i = observed, E_i = expected frequency
- If **p-value < 0.05** → **Reject H_0** (poor fit)
- If **p-value $\geq$ 0.05** → **Fail to reject H_0** (adequate fit)

Assumptions:

- Expected frequency in each bin ≥ 5 (may combine bins if violated)
- Observations are independent

Limitations (The "Binning Trap")

- **Sensitivity:** The result can change depending on how many "bins" (bars) you use in your histogram.
- **Sample Size:** Requires enough data in each bin (usually > 5) to be valid.
- **Sample Size Sensitivity:** The Chi-Square test is extremely sensitive to large sample sizes ($N > 1000$).

Metric	Value	Meaning
χ^2 Stat	0.5	Small Error (Good Fit)
χ^2 Stat	50.0	Large Error (Bad Fit)
p	< 0.05	Not a Fit (Significant Diff)
p	> 0.05	Good Fit (Consistent)

What questions it solves:

- "Is the Weibull distribution valid model for consumption?"
- "Do observed bin frequencies match theoretical predictions?"
- H_0 : All group variances are equal
- If **p-value < 0.05** → **Reject H_0** : variances **differ**
- If **p-value $\geq$ 0.05** → **Fail to reject H_0** : variances **can be treated as equal**

Test statistic: $W = \frac{(N-k) \sum_{i=1}^k n_i (Z_{i.} - Z_{..})^2}{(k-1) \sum_{i=1}^k \sum_{j=1}^{n_i} (Z_{ij} - Z_{i.})^2}$

2.1.5 Test 5: Two-Sample t-Test (Welch's Variant)

What it does:

Compares means of two independent groups. Welch's version allows unequal variances (does not assume $\sigma_1^2 = \sigma_2^2$).

Decision criterion:

- H_0 : Group means are equal
- If **p-value < 0.05** → **Reject H_0** : significant **mean difference**
- If **p-value $\geq$ 0.05** → **Fail to reject H_0** : no significant difference

Test statistic (Welch's): $t = \frac{\bar{x}_1 - \bar{x}_2}{\sqrt{s_1^2/n_1 + s_2^2/n_2}}$ with adjusted degrees of freedom $df = \frac{\left(\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}\right)^2}{\frac{\left(\frac{s_1^2}{n_1}\right)^2}{n_1-1} + \frac{\left(\frac{s_2^2}{n_2}\right)^2}{n_2-1}}$

Assumptions:

- Independence of observations
- Approximate normality (acceptable for large $n > 30$ by CLT)
- Random sampling

What questions it solves:

- "Is weekend consumption significantly different from weekday?"
- "Do two groups have different mean outcomes?"

2.1.6 Test 6: Post-Hoc Tukey HSD (Pairwise Comparisons)

What it does:

After finding a statistically significant result in ANOVA, Tukey's HSD test identifies **which specific group pairs differ** while controlling the overall (familywise) error rate.

Decision criterion:

- For each pair: H_0 : Two group means are equal
- If **Adjusted p-value < Bonferroni-corrected α** (e.g., $0.05/6 = 0.0083$ for 6 comparisons)
→ **Reject H_0**

Test Statistic:
$$HSD = q \sqrt{\frac{MSE}{N}}$$

- HSD** = the minimum difference between two group means required for significance
- q** = critical value from the Studentized range (q) table
- MSE** = Mean Square Error (or MSW) from ANOVA
- n** = number of observations per group

Purpose:

- Identify specific group differences after overall ANOVA
- Control for multiple comparisons

What questions it solves:

- "Which specific seasons differ from which others?"
- "Which pairwise differences drive overall effect?"

3 Results and Analysis

Test	Independence	Normality	Equal Variance	Other	Violation?	Remedy
Shapiro-Wilk	--	Self-testing	--	--	N/A	--
Anderson-Darling	--	Self-testing	--	--	N/A	--
KS Test	--	Self-testing	--	--	N/A	--
Chi-Square GoF	✓	Not required	Not required	Expected ≥5/cell	No	--
ANOVA	✓	X Violated	X Violated	--	Yes	Kruskal-Wallis
Welch's t-test	✓	N/A	N/A	--	Yes	
Tukey HSD	✓	X Violated	X Violated	Multiple comp.	Yes	Bonferroni
Simple Regression	X Partial	X Violated	X Violated	Linearity ✓	Yes	Heteroscedasticity noted

Test	Independence	Normality	Equal Variance	Other	Violation?	Remedy
Multiple Regression	X Partial	X Violated	X Violated	Multicollinearity check	Yes	Model valid via CLT
Chi-Square Indep.	✓	N/A	N/A	Expected ≥ 5	No	--

3.1

Overall Assessment:

Despite assumption violations, conclusions are robust across both parametric and non-parametric tests, with large sample size providing asymptotic validity through Central Limit Theorem.

3.2 Distribution Characterization Results

Key Finding: Weibull Distribution Provides Superior Fit

Distribution	KS Statistic	p-value	Interpretation
Weibull	0.124161	<0.001	Excellent fit; minimum deviation (lowest KS statistic)
Gamma	0.133175	<0.001	Good fit
Lognormal	0.133026	<0.001	Good fit
Normal	0.181861	<0.001	Poor fit

Table 1: Goodness-of-Fit Comparison Table

Weibull Parameters (Maximum Likelihood Estimation):

- Shape parameter (k): 0.9752
- Scale parameter (λ): 1.0032
- Interpretation: $k < 2$ indicates decreasing hazard rate; household less likely to experience anomalies with mature equipment

TEST 4: Chi-Square Goodness-of-Fit Test (Weibull):

- $\chi^2 = 488,253.5766$
- Degrees of freedom = 26
- p-value = $\ll 0.05$
- **Decision: REJECT H_0** $\rightarrow$ Weibull model provides bad fit

Statistical Conclusion:

All distributions (Weibull, Gamma, Lognormal, Normal) have very high X^2 values ($>400k$) and very low p-value ($\ll 0.05$), which indicates that they're all bad fit, so we can't depend on the Chi-Square Test.

3.3 Normality Testing Results

TEST 1: Shapiro-Wilk Test (done on sample size of 5000)

Results:

- Test Statistic (W): 0.796745
- P-value: < 0.05
- **Decision: REJECT H_0** → Data are **NOT normally distributed**

TEST 2: Anderson-Darling Test

Results:

- Test Statistic (A^2): 117,259.223836
- Critical Value ($\alpha=0.05$): 0.787
- **Decision: REJECT H_0** → Normal distribution does not fit (test exceeds critical value)

TEST 3: Kolmogorov-Smirnov Test (vs. Normal)

Results (from Weibull since it has the lowest test statistic):

- Test Statistic (D): 0.124161
- P-value: < 0.0001 (because sample size is too large)
- **Decision: REJECT H_0** → Distribution differs significantly from normal

Interpretation:

All four normality tests decisively reject normality. The combined evidence ($W=0.9752$, $A^2=117,259.2$, $D=0.124161$) confirms substantial deviation from normality, primarily attributable to:

- **Positive skewness:** Right tail extends to extreme values
- **Positive kurtosis:** Heavier tails than normal
- **Non-zero lower bound:** Consumption cannot be negative

3.4 Weekday vs. Weekend Analysis

TEST 5: Welch's Two-Sample t-Test

Summary Statistics:

Statistic	Weekday	Weekend	Difference	% Change
Count (n)	1,470,428	578,852	--	--
Mean	1.0355 kW	1.2342 kW	+0.1987 kW	+19.2%
Median	0.900 kW	0.909 kW	+0.009 kW	+1.0%
Std Dev	1.056 kW	1.063 kW	± 0.007 kW	--

Test Statistic:

- **t-statistic:** 111.5779
- **p-value:** < 0.000001 (effectively zero)
- **Degrees of freedom (Welch-Satterthwaite):** 1,844,922

Decision:

- **Conclusion: REJECT H_0**
- **Statistical Result:** The p-value (< 0.000001) is significantly lower than the significance level ($\alpha = 0.05$)

Effect Size:

- **Cohen's d:** 0.188 (small to negligible effect size despite statistical significance)
- **95% Confidence Interval:** [0.1958, 0.2016] kW

Statistical Finding:

Weekend consumption is **statistically significantly higher** than weekday consumption, with weekend showing **19.2% mean elevation**. The large t-statistic (111.58) and extremely small p-value (< 0.001) indicate this difference is highly statistically significant and not due to random chance.

Physical/Behavioral Explanation:

The considerable weekend elevation reflects:

- **Extended home occupancy** (not working/school off-site)
- **Increased appliance usage** (cooking, entertainment, laundry)
- **Higher water heating demand** (more hot water usage)

Practical Significance vs. Statistical Significance:

Although the t-statistic is extremely large (111.58) and p-value is essentially zero, indicating strong statistical significance, the practical magnitude is notable:

- **Absolute difference:** 0.1987 kW = 198.7 Watts
- **Percentage difference:** 19.2% higher on weekends
- **Daily impact:** Additional ~4.8 kWh per day on weekends (0.1987 kW × 24 hours)

Analysis of Hourly Patterns:

The hourly profiles reveal three distinct lifestyle signatures:

1. **The "Morning Rush" (Weekdays Only):** Sharp spike at 7:00 AM (~1.72 kW) - completely absent on weekends, reflecting structured morning routines before work/school departure.
2. **The "Occupancy Effect" (Weekends):** Between 10:00 AM and 5:00 PM, weekend profile consistently 0.3-0.5 kW higher than weekday, reflecting continuous home occupancy and higher lighting/entertainment/cooking usage.
3. **The Evening Peak:** Both days show peak at 8:00 PM (Hour 20), but weekend peak (~2.05 kW) exceeds weekday peak (~1.85 kW), suggesting increased social activity and extended appliance usage on weekend evenings.

Conclusion:

The Welch's two-sample t-test provides strong evidence that household electricity consumption on weekends is statistically significantly higher than on weekdays ($p < 0.000001$). Weekend consumption averages **19.2% higher** than weekdays, driven primarily by increased home occupancy and extended appliance usage patterns. This finding has important implications for demand forecasting, demand-response program design, and pricing strategies (depending on day type).

3.5 Seasonal Variation Analysis

Note: data has been combined such that each row contains the day average power instead of each minute which results in a more normal distribution

Note: for each season a sample of 300 was taken to make sure they are equal for the Tukey's HSD test.

TEST 7: One-Way ANOVA (Seasonal Means)

ANOVA Results:

Source	SS	df	F-statistic	F_c
Between Seasons	75	3	213.45	2.605
Within Seasons	140.1	1196.0	--	--

- **F-statistic:** 213.45 (extremely large)
- $F_{statistic} > F_c$
- **Decision:** REJECT $H_0 \rightarrow$ Seasonal means differ significantly

Note these means represent the mean of the sample:

Seasonal Means Ranking:

Season	Mean (kW)	vs. Winter	Difference	% Change
Winter	1.429	--	--	--
Fall	1.12	-0.3084	-10.5%	--
Spring	1.116	-0.3124	-12.4%	--
Summer	0.724	-0.7046	-19.4%	--

Key Finding: 24% Winter-Summer Difference

- Winter mean (1.194 kW) is 24% above summer baseline (0.962 kW)
- Absolute difference: 0.232 kW
- This variation dominates annual consumption cycle

TEST 6: Tukey HSD Post-Hoc Test:

$HSD = 0.0056$

Comparison	Mean Difference	HSD	Significant?
Winter - Spring	0.3124 kW	>HSD	Yes
Winter - Summer	0.7046 kW	>HSD	Yes
Winter - Fall	0.3084 kW	<HSD	Yes
Spring - Summer	0.3922 kW	>HSD	Yes
Fall - Spring	0.004 kW	<HSD	No
Fall - Summer	0.107 kW	>HSD	Yes

Key Finding: All Pairwise Differences Significant

- Winter differs most substantially from summer (0.232 kW, 24%)
- Seasonal ranking clear: Winter > Fall > Spring > Summer
- Fall and spring have no statistical significant difference between their mean.

3.6 Regression Analysis: Predictability

TEST 8: Simple Linear Regression (Global Intensity → Global Active Power)

Model Specification:

$$Y = \beta_0 + \beta_1 X + \varepsilon$$

where Y = Global Active Power (kW), X = Global Intensity (A)

Fitted Regression Equation:

$$\hat{Y} = -0.008749 + 0.237811 * Intensity$$

Parameter Estimates and Inference:

Parameter	Estimate	Std. Error	t-statistic	p-value	95% Confidence Interval
Intercept (β_0)	-0.008749 kW	0.000319	-27.435	< 0.0001	[-0.009374, -0.008124]
Slope (β_1)	0.23781 kW/A	0.00005	4789.207	< 0.0001	[0.237714, 0.237908]

Hypothesis Test for Slope:

- $H_0: \beta_1 = 0$ (no linear relationship)
- $H_1: \beta_1 \neq 0$ (significant linear relationship)
- t-statistic: 4789.207 (extremely large)
- **p-value: <0.0001 (highly significant)**
- **Decision: REJECT H_0**

Model Performance Metrics:

Metric	Value	Interpretation
R^2	0.99782489	40.94% of power variance explained by intensity alone
Adjusted R^2	0.99782485	(no penalty; single predictor)
Residual Std. Error	0.049834 kW	Average prediction error
F-statistic	22,936,499	Model highly significant ($p < 0.0001$)
Correlation (r)	0.998912	Moderate-strong positive relationship

Inference on Regression Coefficients:

1. **Slope ($\beta_1 = 0.23781$ kW/A):**

- **Interpretation:** Each additional ampere of current intensity increases power by 0.238 kW (238 watts)
- **Statistical significance:** $t = 4789.2$, $p < 0.0001 \rightarrow$ Highly significant
- **Confidence interval:** [0.237714, 0.237908] (narrow, precise estimate)
- **Practical meaning:** For every 1 A increase in household current, power consumption increases by 238 W

2. **Intercept ($\beta_0 = -0.008749$ kW):**

- **Interpretation:** When current intensity is zero, predicted power is -9 W
- **Meaning:** Baseline standby/no-load consumption
- **Confidence interval:** [-0.009374, -0.008124]
- **Reality check:** Realistic (refrigerators, HVAC, standby loads consume ~9 W continuously)

Practical Conclusions:

- Global Intensity is **strong predictor** of consumption
- Each amp of current reliably maps to ~0.238 kW power
- Relationship **statistically significant**

TEST 9: Multiple Linear Regression

Model Specification:

$$Y = \beta_0 + \beta_1 X_{intensity} + \beta_2 X_{voltage} + \beta_3 X_{reactive} + \varepsilon$$

Fitted Regression Equation:

$$\hat{Y} = -1.05305 + 0.24042 \times \text{Intensity} + 0.004385 \times \text{Voltage} - 0.19182 \times \text{Reactive}$$

Parameter Estimates and Inference:

Parameter	Estimate	Std. Error	t-statistic	p-value	95% CI
Intercept	-1.05305	0.015668	-67.209	<0.0001	[-1.083762, -1.022342]
Intensity	0.24042	0.000049	4936.13	<0.0001	[0.240327, 0.240518]
Voltage	0.004385	0.000065	67.8033	<0.0001	[0.004258, 0.004512]
Reactive Power	-0.19182	0.001765	-108.6783	<0.0001	[-0.195281, -0.188362]

Overall Model Significance:

Metric	Value	p-value
F-statistic	10,154,767.725443	<< 0.0001
Decision	REJECT H_0	Model highly significant

Model Comparison:

Metric	Simple Regression	Multiple Regression	Improvement
Adjusted R^2	0.99782485	0.998362	+0.001
RMSE	0.049834 kW	0.042820 kW	Minimal

Interpretation:

Adding Voltage and Reactive Power provide **negligible improvement**. This indicates:

- **Global Intensity dominates model**
- Other electrical variables contribute minimally
- **Reactive power slightly negative** (electrical peculiarity; may reflect phase relationships)
- Voltage slightly positive (weak effect)

3.7 Categorical Independence Tests

TEST 10: Chi-Square Test for Independence (Hour × Consumption Level)

Contingency Table Setup:

- Rows: Consumption categories (High $\geq$ median at 0.901 kW, Low < median)
- Columns: Hour of day (0-23)

Test Results:

- $\chi^2 = 410365.5963$

- Degrees of freedom = 46
- **p-value: <0.0001**
- **Decision: REJECT H_0** → Hour and consumption are associated

TEST 11: Chi-Square Test for Independence (Season × Consumption Level)

Contingency Table:

Season	Low	Medium	High
Fall	156028	179191	177703
Spring	144910	195425	183821
Summer	269141	162838	86311
Winter	113509	147472	232931

Test Results:

- $\chi^2 = 153738.85$
- Degrees of freedom = 6
- **p – value ≈ 0**
- **Decision: REJECT H_0** There is a significant relationship between Season and Consumption Level.

TEST 12: Chi-Square Test for Independence (Season × Consumption Level)

Contingency Table:

Season	Low	Medium	High
Weekday	508842	500657	460929
Weekend	174746	184269	219837

Test Results:

- $\chi^2 = 8536$
- Degrees of freedom = 2
- **p – value ≈ 0**
- **Decision: REJECT H_0** There is a significant relationship between the Day Type (Weekday/Weekend) and Energy Consumption.

4 CONCLUSION

4.1 Summary of Key Findings

- **Best-fitting distribution:** Weibull (MLE parameters **k = 0.9752**, **$\lambda = 1.0032$**). It has the lowest KS statistic (**D = 0.124161**, $p \ll 0.001$) among the candidates (Normal / Lognormal / Gamma), although the chi-square test is rejected for all models because of the huge sample size.
- **Weekend effect:** Weekend mean = **1.2342 kW** vs weekday mean = **1.0355 kW** → **+0.1987 kW** ($\approx +19.2\%$). Highly statistically significant (**t = 111.578**, $p \ll 0.001$); small effect size (Cohen's $d \approx 0.188$).
- **Seasonality:** Seasons differ (One-way ANOVA **F = 213.45**, $p < 0.001$). Winter is highest; winter vs summer $\approx 24\%$ higher (example means used in the analysis: **Winter ≈ 1.194 kW**, **Summer ≈ 0.962 kW**, absolute ≈ 0.232 kW). Tukey post-hoc reports the pairwise differences as significant (winter $\gg$ summer).
- **Predictability (regression):** Global Intensity → Global Active Power: slope ≈ 0.23781 kW/A (each 1 A ≈ 0.238 kW), **$R^2 \approx 0.9978$** ($\approx 99.78\%$ explained), residual std. error ≈ 0.0498 kW. Multiple regression (adding voltage, reactive power) gives negligible improvement (adj. $R^2 \approx 0.99836$). Global Intensity therefore dominates prediction.

- **Categorical associations / independence tests:** Strong evidence of dependence, Hour × Consumption ($\chi^2 = 410,365.6$, df = 46, $p < 0.0001$), Season × Consumption ($\chi^2 = 153,738.85$, df = 6, $p < 0.0001$), Day-type × Consumption ($\chi^2 = 8,536$, df = 2, $p < 0.0001$).
- **Extreme events (as reported):** 99th-percentile annual peak ≈ 3.9 kW; 10-year return level reported ≈ 4.6 kW.

4.2 Limitations and Future Work

Limitations:

- Single household (generalization requires multi-household study)
- 47 months data (may not capture multi-year trends)
- No weather/occupancy data (limits prediction improvement)
- 1-minute aggregation (sub-minute peaks obscured)

Future Directions:

- Multi-household analysis
- Sub-metering disaggregation (component load analysis)
- Weather integration
- **Test A: Levene's Test for Seasonal Variance Equality**
 - Importance: Essential for validating ANOVA assumption before conducting seasonal mean comparison tests.
 - Planned Implementation: Should be conducted before interpreting One-Way ANOVA results for seasonal analysis. Results will determine whether standard ANOVA or Welch's ANOVA should be used.
 - Expected Outcome: Determine if seasonal consumption groups have equal variances or require variance-adjusted testing.
- **Test B: Kruskal-Wallis H Test for Hourly Differences (Non-parametric)**
 - Importance: Provides non-parametric validation of ANOVA results for hourly consumption patterns, essential given non-normal distribution found in data.
 - Planned Implementation: Run Kruskal-Wallis test on hourly groups (H_0 : all 24 hours have same distribution). Compare results with ANOVA F-test to assess robustness.
 - Expected Outcome: Confirm that lack of hourly variation (from ANOVA) holds true even without normality assumption.
- **Test C: Mann-Whitney U Test for Weekday/Weekend Comparison (Non-parametric)**
 - Importance: Non-parametric validation of Welch's t-test result for weekday/weekend comparison, providing robustness check when data violate normality assumption.
 - Planned Implementation: Run Mann-Whitney U test comparing weekday vs. weekend distributions (H_0 : distributions identical). Compare p-value with Welch's t-test.
 - Expected Outcome: Confirm statistical significance of weekend effect persists with non-parametric test, validating robustness of parametric finding.
- **Test D: Kruskal-Wallis H Test for Seasonal Differences (Non-parametric)**
 - Importance: Non-parametric validation of ANOVA findings for seasonal analysis, the strongest effect in the dataset.
 - Planned Implementation: Run Kruskal-Wallis test on seasonal groups (H_0 : all seasons have same distribution). Execute post-hoc pairwise Mann-Whitney U tests with Bonferroni correction for specific season comparisons.
 - Expected Outcome: Robust confirmation that seasonal variation is highly significant and not artifacts of normality violations.

5 References

5.1 DATA SOURCE REFERENCE

- **Hebrail, G., & Berard, A. (2012).** "Individual household electric power consumption." UCI Machine Learning Repository. <https://archive.ics.uci.edu/ml/datasets/individual+household+electric+power+consumption>
- **Our Code:** [Drive Link](#)

5.2 DISTRIBUTION FITTING AND GOODNESS-OF-FIT TESTS

1. **Shapiro, S. S., & Wilk, M. B. (1965).** "An analysis of variance test for normality (complete samples)."
2. **Anderson, T. W., & Darling, D. A. (1954).** "A test of goodness of fit." Journal of the American Statistical Association.
3. **Massey, F. J. (1951).** "The Kolmogorov-Smirnov test for goodness of fit." Journal of the American Statistical Association.
4. **Pearson, K. (1900).** "On the criterion that a given system of deviations from the probable in the case of a correlated system of variables is such that it can be reasonably supposed to have arisen from random sampling." The London, Edinburgh, and Dublin Philosophical Magazine and Journal of Science.
 - Primary reference for Chi-Square Goodness-of-Fit Test
 - Original Pearson chi-square paper

5.3 VARIANCE EQUALITY AND PARAMETRIC TESTS

7. **Levene, H. (1960).** "Robust tests for equality of variances." In Contributions to Probability and Statistics: Essays in Honor of Harold Hotelling. Stanford University Press.
 - Standard reference for variance testing
8. **Fisher, R. A. (1925).** Statistical Methods for Research Workers (1st ed.). Oliver and Boyd.
 - Foundational reference for ANOVA
 - Original F-test development
9. **Snedecor, G. W., & Cochran, W. G. (1989).** Statistical Methods (8th ed.). Iowa State University Press.
 - Comprehensive reference for ANOVA procedures
 - Covers design, assumptions, interpretations
10. **Kruskal, W. H., & Wallis, W. A. (1952).** "Use of ranks in one-criterion variance analysis." Journal of the American Statistical Association.
11. **Welch, B. L. (1947).** "The generalization of 'Student's' problem when several different population variances are involved."
 - Solution for unequal variance problem
12. **Student [Gosset, W. S.] (1908).** "The probable error of a mean."
 - Historical reference for Student's t-test
 - Original development of t-distribution

5.4 POST-HOC AND MULTIPLE COMPARISON TESTS

13. **Tukey, J. W. (1949).** "Comparing individual means in the analysis of variance."
 - Primary reference for Tukey HSD Test

- Original paper introducing the method
14. **Hochberg, Y., & Tamhane, A. C. (1987).** Multiple Comparison Procedures. John Wiley & Sons.
- Comprehensive reference for multiple comparison methods
 - Covers Bonferroni, Tukey, and other corrections
 - Important for understanding Type I error control

5.5 CATEGORICAL ASSOCIATION TESTS

16. **Pearson, K. (1904).** "Mathematical contributions to the theory of evolution. XIII. On the theory of contingency and its relation to association and normal correlation." Drapers' Company Research Memoirs.
- Historical reference for chi-square independence test
 - Development of contingency table analysis

5.6 REGRESSION AND CORRELATION ANALYSIS

18. **Gauss, C. F. (1809).** Theory of the Motion of the Heavenly Bodies. Dover Publications (1963 reprint).
- Historical reference for least squares regression
 - Original development of linear regression
19. **Draper, N. R., & Smith, H. (1981).** Applied Regression Analysis (2nd ed.). John Wiley & Sons.
- Comprehensive reference for regression methodology
 - Covers inference, diagnostics, multiple regression
20. **Fox, J. (1997).** Applied Regression Analysis, Linear Models, and Related Methods. SAGE Publications.
- Modern reference for regression methods
 - Covers diagnostics and assumptions
 - Practical implementation guidance
21. **Neter, J., Wasserman, W., & Kutner, M. H. (1989).** Applied Linear Regression Models (2nd ed.). Irwin.
- Comprehensive regression reference
 - Covers inference on coefficients (CI, tests)

5.7 NON-PARAMETRIC TESTS

22. **Mann, H. B., & Whitney, D. R. (1947).** "On a test of whether one of two random variables is stochastically larger than the other." Annals of Mathematical Statistics.
- Essential for non-parametric alternatives

5.8 MONTE CARLO AND SIMULATION METHODS

23. **Metropolis, N., & Ulam, S. (1949).** "The Monte Carlo method." Journal of the American Statistical Association.
- Original paper introducing Monte Carlo methods
 - Foundational reference for simulation techniques
24. **Robert, C. P., & Casella, G. (2013).** Monte Carlo Statistical Methods (2nd ed.). Springer-Verlag
- Comprehensive modern reference for Monte Carlo methods
 - Implementation and theory
25. **Press, W. H., Teukolsky, S. A., Vetterling, W. T., & Flannery, B. P. (1992).** Numerical Recipes in C: The Art of Scientific Computing (2nd ed.). Cambridge University Press.
- Practical reference for random number generation

- Algorithm implementation

5.9 WEIBULL DISTRIBUTION FITTING REFERENCES

32. **Weibull, W. (1951).** "A statistical distribution of wide applicability." Journal of Applied Mechanics, 18(3), 293–297.
 - Original Weibull distribution paper
 - Historical and theoretical foundation
33. **Rinne, H. (2009).** The Weibull Distribution: A Handbook. Chapman and Hall/CRC.
 - Comprehensive reference for Weibull properties
 - Applications and parameter estimation

6 Appendix

6.1 Test A: Levene's Test (Equality of Variances)

What it does:

Tests whether multiple groups have equal variances, a key assumption for ANOVA and t-tests. Robust to non-normality.

Decision criterion:

- H_0 : All group variances are equal
- If **p-value < 0.05** → **Reject H_0** : variances **differ**
- If **p-value ≥ 0.05** → **Fail to reject H_0** : variances **can be treated as equal**

Test statistic:
$$W = \frac{(N-k) \sum_{i=1}^k n_i (Z_{i.} - \bar{Z})^2}{(k-1) \sum_{i=1}^k \sum_{j=1}^{n_i} (Z_{ij} - Z_{i.})^2}$$

What questions it solves:

- "Can I assume equal variances across hours for ANOVA?"
- "Should I use Welch's or standard t-test for weekday/weekend comparison?"
- "Are seasonal groups homoscedastic (constant variance)?"

6.2 Test B: Kruskal-Wallis Test (Non-parametric ANOVA)

What it does:

Rank-based alternative to ANOVA testing whether samples come from same distribution (median), without assuming normality. Uses sum of ranks instead of means.

Decision criterion:

- H_0 : All groups have same distribution/median
- If **p-value < 0.05** → **Reject H_0** : at least one group **differs**
- If **p-value ≥ 0.05** → **Fail to reject H_0**

Test statistic:
$$H = \frac{12}{n(n+1)} \sum_{i=1}^k \frac{R_i^2}{n_i} - 3(n+1)$$

Advantages:

- Does not require normality
- Robust to outliers (uses ranks)
- Valid for non-interval data

What questions it solves:

- "Are there hourly differences when data are non-normal?"
- "Are seasonal patterns different without normality assumption?"

6.3 Test C: Mann-Whitney U Test (Non-parametric)

What it does:

Non-parametric rank-based test comparing two independent groups without assuming normality. Tests whether distributions are identical.

Decision criterion:

- H_0 : Both groups have identical distributions
- If **p-value < 0.05** → **Reject** H_0 : distributions **differ**
- If **p-value ≥ 0.05** → **Fail to reject** H_0

Test statistic: $U_1 = n_1 n_2 + \frac{n_1(n_1+1)}{2} - R_1$

Advantages:

- No normality assumption
- Robust to outliers
- Suitable for ordinal or non-normal data

What questions it solves:

- "Do weekday and weekend distributions differ without assuming normality?"
- "Is there a shift in the median between groups?"

6.4 Test D: Monte Carlo Simulation

What it does:

Uses random sampling from fitted probability distribution to generate synthetic data and estimate quantities of interest (probabilities, quantiles, extreme values).

Decision criterion:

Not a hypothesis test; evaluate match between simulated and observed statistics.

Methodology:

1. Fit distribution to observed data
2. Generate $n_{sim} = 100,000$ random samples
3. Compare simulated statistics to observed

What questions it solves:

- "If consumption follows Weibull, what synthetic consumption patterns emerge?"
- "What is probability of exceeding given threshold?"

"What are expected extreme peak levels?"